


Seasonal time series

ARIMA MODELS IN PYTHON




James Fulton

Climate informatics researcher

Seasonal data

- Has predictable and repeated patterns
- Repeats after any amount of time

Seasonal decomposition



time series = trend + seasonal + residual

Seasonal decomposition using statsmodels

```
# Import  
from statsmodels.tsa.seasonal import seasonal_decompose
```


```
# Decompose data  
decomp_results = seasonal_decompose(df['IPG3113N'], period=12)
```

```
type(decomp_results)
```


```
statsmodels.tsa.seasonal.DecomposeResult
```

Seasonal decomposition using statsmodels


```
# Plot decomposed data  
decomp_results.plot()  
plt.show()
```



Finding seasonal period using ACF

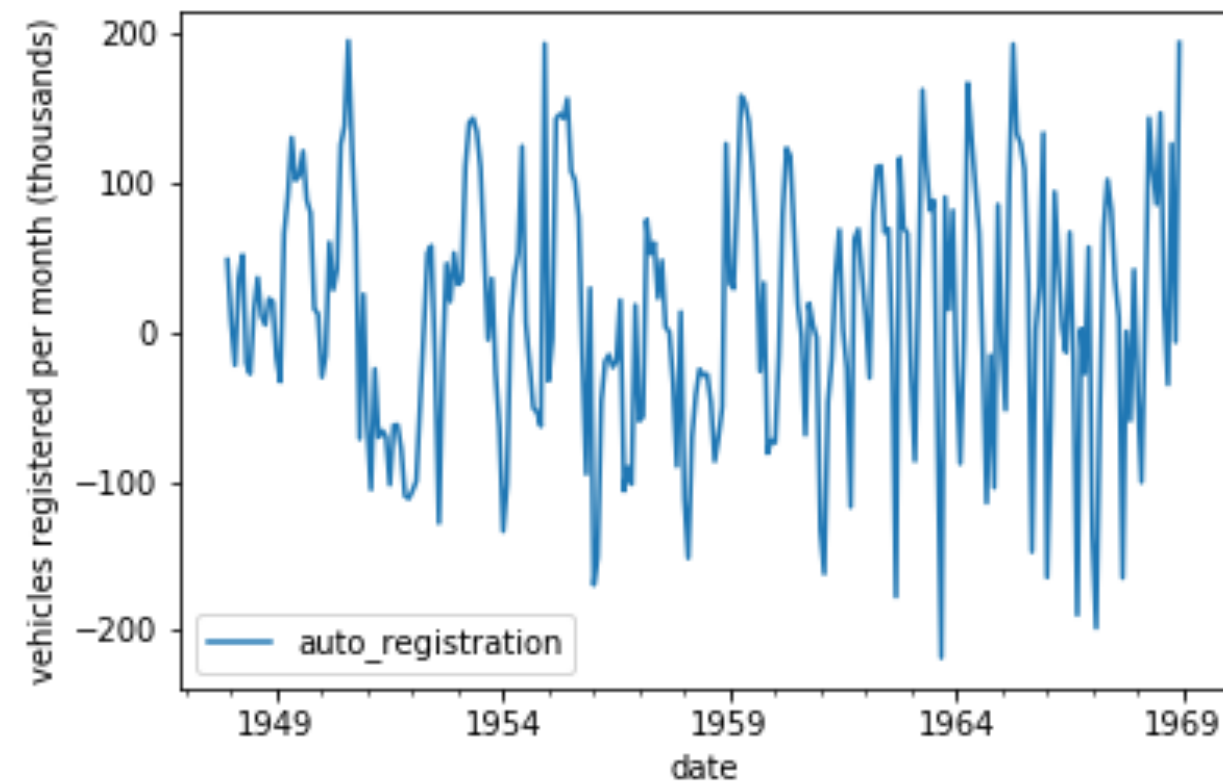


Identifying seasonal data using ACF



Detrending time series


```
# Subtract long rolling average over N steps  
df = df - df.rolling(N).mean()  
  
# Drop NaN values  
df = df.dropna()
```




Identifying seasonal data using ACF

```
# Create figure
fig, ax = plt.subplots(1,1, figsize=(8,4))

# Plot ACF
plot_acf(df.dropna(), ax=ax, lags=25, zero=False)
plt.show()
```



ARIMA models and seasonal data



The SARIMA model

Seasonal ARIMA = SARIMA

- Non-seasonal orders
 - p : autoregressive order
 - d : differencing order
 - q : moving average order

$\text{SARIMA}(p,d,q)(P,D,Q)_S$

- Seasonal Orders
 - P : seasonal autoregressive order
 - D : seasonal differencing order
 - Q : seasonal moving average order
 - S : number of time steps per cycle

The SARIMA model

ARIMA(2,0,1) model :

$$y_t = a_1 y_{t-1} + a_2 y_{t-2} + m_1 \epsilon_{t-1} + \epsilon_t$$

SARIMA(0,0,0)(2,0,1)₇ model:

$$y_t = a_7 y_{t-7} + a_{14} y_{t-14} + m_7 \epsilon_{t-7} + \epsilon_t$$

Fitting a SARIMA model

```
# Imports
statsmodels.tsa.statespace.sarimax import SARIMAX

# Instantiate model
model = SARIMAX(df, order=(p,d,q), seasonal_order=(P,D,Q,S))

# Fit model
results = model.fit()
```


Seasonal differencing

Subtract the time series value of one season ago

$$\Delta y_t = y_t - y_{t-S}$$

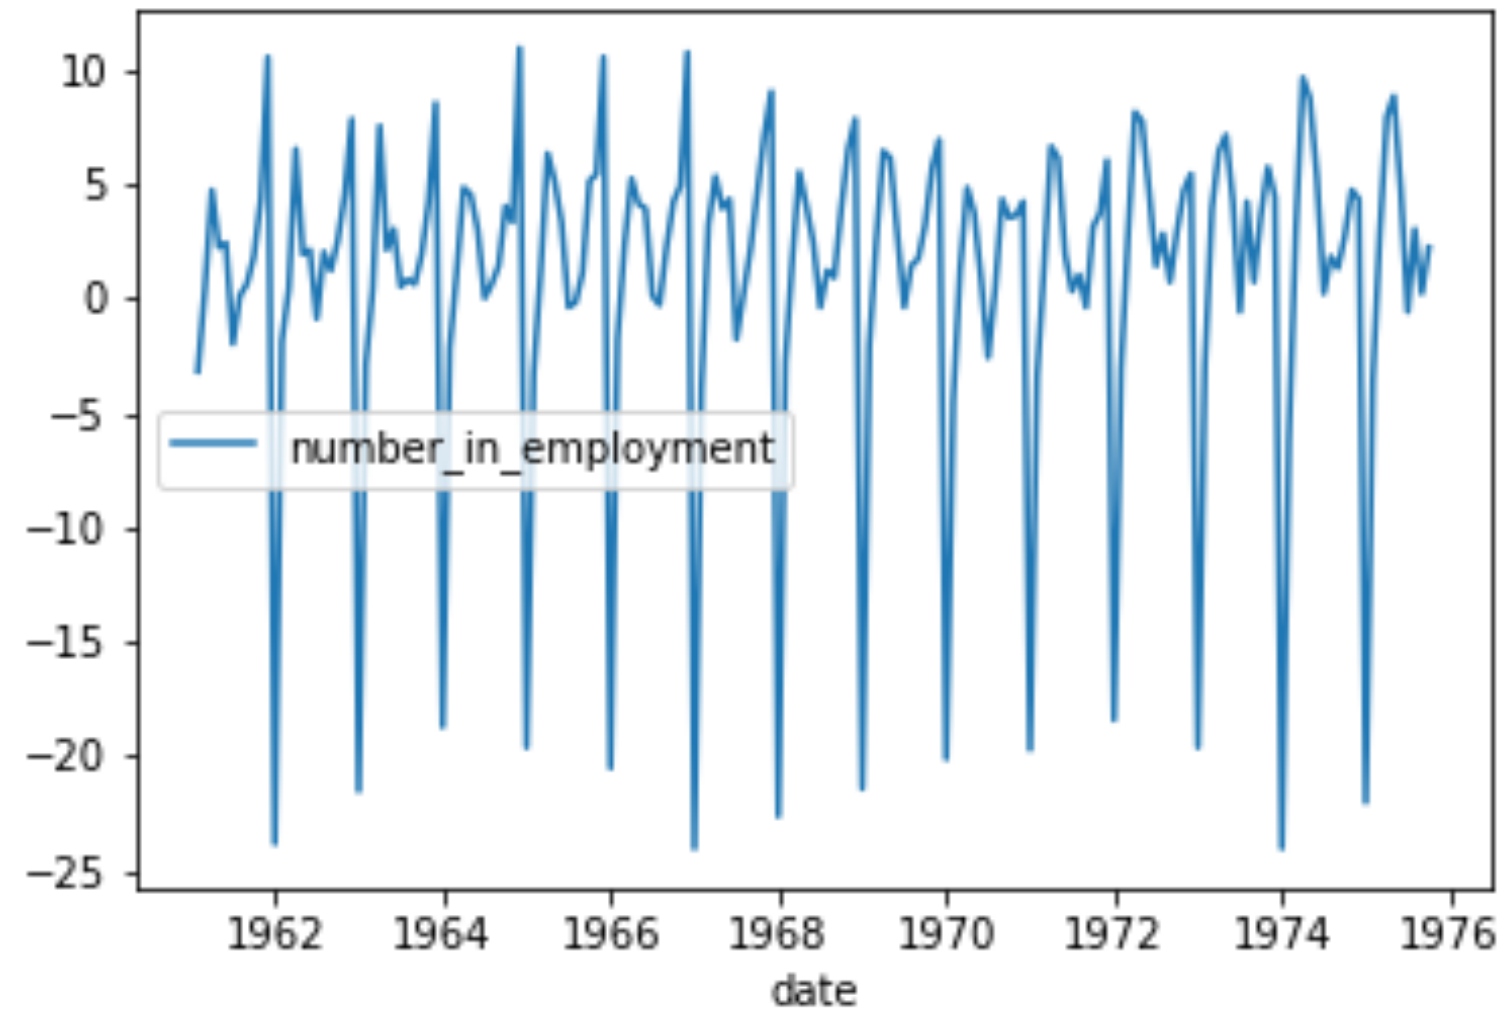
```
# Take the seasonal difference  
df_diff = df.diff(S)
```

Differencing for SARIMA models



Time series

Differencing for SARIMA models




First difference of time series

Differencing for SARIMA models



First difference and first seasonal difference of time series

Finding p and q



Plotting seasonal ACF and PACF


```
# Create figure
fig, (ax1, ax2) = plt.subplots(2,1)

# Plot seasonal ACF
plot_acf(df_diff, lags=[12,24,36,48,60,72], ax=ax1)

# Plot seasonal PACF
plot_pacf(df_diff, lags=[12,24,36,48,60,72], ax=ax2)

plt.show()
```

Finding P and Q



Searching over model orders

```
import pmdarima as pm
```

```
results = pm.auto_arima(df)
```

```
Performing stepwise search to minimize aic
```

```
ARIMA(2,0,2)(1,1,1)[12] intercept : AIC=inf, Time=3.33 sec
```

```
ARIMA(0,0,0)(0,1,0)[12] intercept : AIC=2648.467, Time=0.062 sec
```

```
ARIMA(1,0,0)(1,1,0)[12] intercept : AIC=2279.986, Time=1.171 sec
```

```
...
```

```
ARIMA(3,0,3)(1,1,1)[12] intercept : AIC=2173.508, Time=12.487 sec
```

```
ARIMA(3,0,3)(0,1,0)[12] intercept : AIC=2297.305, Time=2.087 sec
```

```
Best model: ARIMA(3,0,3)(1,1,1)[12]
```

```
Total fit time: 245.812 seconds
```


pmdarima results

```
print(results.summary())
```

```
=====
                        Statespace Model Results
=====
Dep. Variable:          real values    No. Observations:          300
Model:                  SARIMAX(2, 0, 0)  Log Likelihood            -408.078
Date:                   Tue, 28 May 2019  AIC                        822.156
Time:                   15:53:07         BIC                        833.267
Sample:                 01-01-2013      HQIC                       826.603
                        - 10-27-2013
Covariance Type:        opg
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
ar.L1         0.2189     0.054      4.072     0.000     0.114     0.324
ar.L2         0.1960     0.054      3.626     0.000     0.090     0.302
sigma2        0.8888     0.073     12.160     0.000     0.746     1.032
=====
Ljung-Box (Q):                32.10    Jarque-Bera (JB):                0.02
Prob(Q):                      0.81    Prob(JB):                      0.99
Heteroskedasticity (H):        1.28    Skew:                          -0.02
Prob(H) (two-sided):           0.21    Kurtosis:                      2.98
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
results.plot_diagnostics()
```



Non-seasonal search parameters

```
results = pm.auto_arima( df,                # data
                        d=0,                # non-seasonal difference order
                        start_p=1,          # initial guess for p
                        start_q=1,          # initial guess for q
                        max_p=3,            # max value of p to test
                        max_q=3,            # max value of q to test
                        )
```

¹ https://www.alkaline-ml.com/pmdarima/modules/generated/pmdarima.arima.auto_arima.html

Seasonal search parameters

```
results = pm.auto_arima(df,          # data
                        ... ,        # non-seasonal arguments
                        seasonal=True, # is the time series seasonal
                        m=7,          # the seasonal period
                        D=1,          # seasonal difference order
                        start_P=1,    # initial guess for P
                        start_Q=1,    # initial guess for Q
                        max_P=2,      # max value of P to test
                        max_Q=2,      # max value of Q to test
                        )
```


Other parameters

```
results = pm.auto_arima(df,                # data
                        ...,                # model order parameters
                        information_criterion='aic', # used to select best model
                        trace=True,          # print results whilst training
                        error_action='ignore', # ignore orders that don't work
                        stepwise=True,       # apply intelligent order search
                        )
```

Saving model objects

```
# Import
```

```
import joblib
```

```
# Select a filepath
```

```
filepath = 'localpath/great_model.pkl'
```

```
# Save model to filepath
```

```
joblib.dump(model_results_object, filepath)
```

Saving model objects


```
# Select a filepath
filepath = 'localpath/great_model.pkl'

# Load model object from filepath
model_results_object = joblib.load(filepath)
```

Updating model

```
# Add new observations and update parameters  
model_results_object.update(df_new)
```

Update comparison




Let's practice!

ARIMA MODELS IN PYTHON

SARIMA and Box-Jenkins

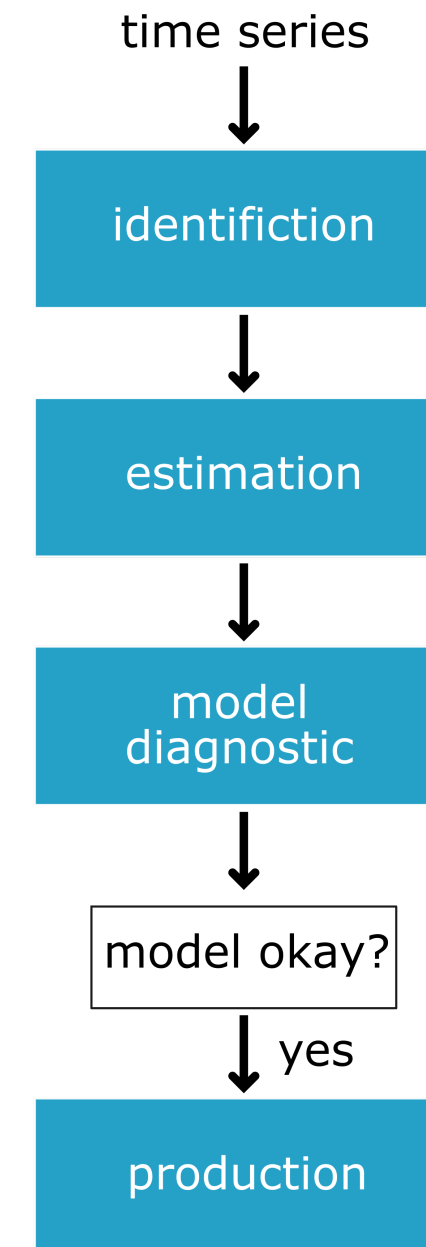
ARIMA MODELS IN PYTHON



James Fulton

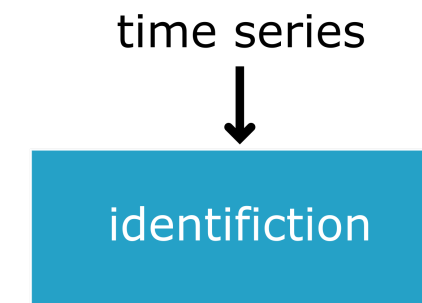
Climate informatics researcher

Box-Jenkins




Box-Jenkins with seasonal data

- Determine if time series is seasonal
- Find seasonal period
- Find transforms to make data stationary
 - Seasonal and non-seasonal differencing
 - Other transforms




Mixed differencing


- D should be 0 or 1
- $d + D$ should be 0-2



Weak vs strong seasonality




- Weak seasonal pattern
- Use seasonal differencing if necessary




- Strong seasonal pattern
- Always use seasonal differencing

Additive vs multiplicative seasonality




- Additive series = trend + season
- Proceed as usual with differencing



- multiplicative series = trend x season
- Apply log transform first - `np.log`

Multiplicative to additive seasonality



The SARIMAX model

S - seasonal
A - autoregressive
R - integrated
I - moving average
M - exogenous
A -
X -

Time series modeling framework

- Test for stationarity and seasonality
- Find promising model orders
- Fit models and narrow selection with AIC/BIC
- Perform model diagnostics tests
- Make forecasts
- Save and update models

